



ÉCOLE POLYTECHNIQUE
FÉDÉRALE DE LAUSANNE

Data-Parallel Programming

Parallel Programming in Scala

Aleksandar Prokopec

Data-Parallelism

Previously, we learned about task-parallel programming.

A form of parallelization that distributes execution processes across computing nodes.

We know how to express parallel programs with task and parallel constructs.

Data-Parallelism

Previously, we learned about task-parallel programming.

A form of parallelization that distributes execution processes across computing nodes.

We know how to express parallel programs with task and parallel constructs.

Next, we learn about the data-parallel programming.

A form of parallelization that distributes data across computing nodes.

Data-Parallel Programming Model

The simplest form of data-parallel programming is the parallel for loop.

Example: initializing the array values.

Data-Parallel Programming Model

The simplest form of data-parallel programming is the parallel for loop.

Example: initializing the array values.

```
def initializeArray(xs: Array[Int])(v: Int): Unit
```

Data-Parallel Programming Model

The simplest form of data-parallel programming is the parallel for loop.

Example: initializing the array values.

```
def initializeArray(xs: Array[Int])(f: Int) => Unit = {  
  for (i <- (0 until xs.length).par) {  
    f(i)  
  }  
}
```

Data-Parallel Programming Model

The simplest form of data-parallel programming is the parallel for loop.

Example: initializing the array values.

```
def initializeArray(xs: Array[Int])(v: Int): Unit = {  
  for (i <- (0 until xs.length).par) {  
    xs(i) = v  
  }  
}
```

Data-Parallel Programming Model

The simplest form of data-parallel programming is the parallel for loop.

Example: initializing the array values.

```
def initializeArray(xs: Array[Int])(v: Int): Unit = {  
  for (i <- (0 until xs.length).par) {  
    xs(i) = v  
  }  
}
```

The parallel for loop is not functional – it can only affect the program through side-effects.

Data-Parallel Programming Model

The simplest form of data-parallel programming is the parallel for loop.

Example: initializing the array values.

```
def initializeArray(xs: Array[Int])(v: Int): Unit = {  
  for (i <- (0 until xs.length).par) {  
    xs(i) = v  
  }  
}
```

The parallel for loop is not functional – it can only affect the program through side-effects.

As long as iterations of the parallel loop write to separate memory locations, the program is correct.

Example: Mandelbrot Set


Although simple, parallel for loop allows writing interesting programs.

Render a set of complex numbers in the plane for which the sequence $z_{n+1} = z_n^2 + c$ does not approach infinity.

Example: Mandelbrot Set

Although simple, parallel for loop allows writing interesting programs.

Render a set of complex numbers in the plane for which the sequence $z_{n+1} = z_n^2 + c$ does not approach infinity.



Example: Mandelbrot Set

We approximate the definition of the Mandelbrot set – as long as the absolute value of z_n is less than 2, we compute z_{n+1} until we do `maxIterations`.

```
private def computePixel(xc: Double, yc: Double, maxIterations: Int): Int = {  
  var i = 0  
  var x, y = 0.0  
  while (x * x + y * y < 4 && i < maxIterations) {  
    val xt = x * x - y * y + xc  
    val yt = 2 * x * y + yc  
    x = xt; y = yt  
    i += 1  
  }  
  color(i)  
}
```

Example: Mandelbrot Set (Data-Parallel)

How do we render the set using data-parallel programming?

```
def parRender(): Unit = {  
  for (idx <- (0 until image.length).par) {  
    val (xc, yc) = coordinatesFor(idx)  
    image(idx) = computePixel(xc, yc, maxIterations)  
  }  
}
```

Rendering the Mandelbrot Set: Demo

Time for a demo!

Rendering the Mandelbrot Set: Demo

Time for a demo!

Summary:

- ▶ task-parallel implementation – the slowest.
- ▶ data-parallel implementation – about $2\times$ faster.


Workload

Different data-parallel programs have different workloads.

Workload is a function that maps each input element to the amount of work required to process it.


Uniform Workload

Defined by a constant function: $w(i) = \text{const}$



Uniform Workload


Defined by a constant function: $w(i) = \text{const}$



Easy to parallelize.


Irregular Workload

Defined by an arbitrary function: $w(i) = f(i)$



Irregular Workload

Defined by an arbitrary function: $w(i) = f(i)$




In the Mandelbrot case: $w(i) = \#iterations$

The workload depends on the problem instance.

Irregular Workload

Defined by an arbitrary function: $w(i) = f(i)$



In the Mandelbrot case: $w(i) = \#iterations$

The workload depends on the problem instance.

Goal of the *data-parallel scheduler*: efficiently balance the workload across processors without any knowledge about the $w(i)$.